

```
In [1]: from google.colab import drive
drive.mount('/content/drive')
```

Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect_uri=urn%3Aietf%3Awg%3Aoauth%3A2.0%3Aoob&scope=email%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fdocs.test%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fdrive%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fdrive.photos.readonly%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fpeopleapi.readonly&response_type=code

Enter your authorization code:
.....

Mounted at /content/drive

In this task, I use 15-Scene to train VGG16 from-scratch and pre-trained. And I found that from-scratch is not good even drop in local-minimum situation, I try to run more times and change loss function and learning rate but still not work well, but for pre-trained is good, the accuracy is 86% for one time 85% for three times

VGG16---15-Scene dataset

import necessary package and vgg16

```
In [0]: import tensorflow as tf
import numpy as np
import keras
import matplotlib.pyplot as plt
from keras.models import Model
from keras.layers import Dense, Flatten, Dropout
import cv2
from keras.applications.vgg16 import VGG16
from keras.optimizers import SGD
from tqdm import tqdm_notebook as tqdm
import os
from keras import optimizers
from keras.callbacks import EarlyStopping
```

process dataset and split it to training dataset, testing dataset

```

In [0]: img_size = 224
        num_classes = 15
        maxepoches = 30
        batch_size = 50

        DATASET_PATH = '/content/drive/My Drive/Colab Notebooks/VGG16_Practice/15-S
        one_hot_lookup = np.eye(num_classes)
        dataset_x = []
        dataset_y = []
        for category in sorted(os.listdir(DATASET_PATH)):
            for fname in os.listdir(DATASET_PATH+"/"+category):
                img = cv2.imread(DATASET_PATH+"/"+category+"/"+fname, 2)
                img = cv2.resize(img, (img_size,img_size))
                dataset_x.append(np.reshape(img, [img_size,img_size,1]))
                dataset_y.append(np.reshape(one_hot_lookup[int(category)], [num_cla

        # train_y = keras.utils.to_categorical(train_y).astype('float32')
        # test_y = keras.utils.to_categorical(test_y).astype('float32')

```

```

In [0]: dataset_x = np.array(dataset_x)
        dataset_y = np.array(dataset_y)

```

```

In [0]: """shuffle dataset"""
        p = np.random.permutation(len(dataset_x))
        dataset_x = dataset_x[p]
        dataset_y = dataset_y[p]

        test_x = dataset_x[:int(len(dataset_x)/10)]
        test_y = dataset_y[:int(len(dataset_x)/10)]
        train_x = dataset_x[int(len(dataset_x)/10):]
        train_y = dataset_y[int(len(dataset_x)/10):]

```

print train_x train_y test_x test_y shape , founding those part is one channel beacuse they are gray image

```

In [7]: print(train_x.shape)
        print(train_y.shape)
        print(test_x.shape)
        print(test_y.shape)

```

```

(4037, 224, 224, 1)
(4037, 15)
(448, 224, 224, 1)
(448, 15)

```

In order fit with VGG16 network , I convert one channel to three channel

```
In [8]: train_x = [cv2.cvtColor(cv2.resize(i,(224,224)), cv2.COLOR_GRAY2BGR) for i
train_x = np.concatenate([arr[np.newaxis] for arr in train_x]).astype('float32')

test_x = [cv2.cvtColor(cv2.resize(i,(224,224)), cv2.COLOR_GRAY2BGR) for i
test_x = np.concatenate([arr[np.newaxis] for arr in test_x]).astype('float32')

print("train_x shape :",train_x.shape)
print("test_x shape:",test_x.shape)

train_x shape : (4037, 224, 224, 3)
test_x shape: (448, 224, 224, 3)
```

```
In [0]: train_x /= 255
test_x /= 255
```

VGG16 from Scratch

This part is VGG16 from scrtech : weights=None and layer.trainable=True

```
In [10]: model_vgg =VGG16(include_top=False,input_shape=(224,224,3),weights=None)
for layer in model_vgg.layers:
    layer.trainable=True
#model=Flatten(name='Flatten')(model_vgg.output)
x = model_vgg.get_layer('block5_pool').output
x = Flatten(name='flatten')(x)
x = Dense(512, activation='relu', name='fc1')(x)
x = Dense(4096, activation='relu', name='fc2')(x)
x = Dense(15, activation='softmax', name='predictions')(x)
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:66: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:541: The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:4432: The name tf.random_uniform is deprecated. Please use tf.random.uniform instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:4267: The name tf.nn.max_pool is deprecated. Please use tf.nn.max_pool2d instead.

```
In [11]: #model=Dense(15,activation='softmax')(model)
model_vgg_scene=Model(inputs=model_vgg.input, outputs=x)
model_vgg_scene.summary()
```

Model: "model_1"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	(None, 224, 224, 3)	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
flatten (Flatten)	(None, 25088)	0
fc1 (Dense)	(None, 512)	12845568
fc2 (Dense)	(None, 4096)	2101248
predictions (Dense)	(None, 15)	61455
=====		
Total params: 29,722,959		
Trainable params: 29,722,959		

Non-trainable params: 0

```
In [0]: def learning_rate_schedule(epoch):
        if epoch <= 10:
            return 1e-4 # 0.00001
        elif epoch <= 20:
            return 1e-5
        elif epoch <= 30:
            return 1e-6
        else:
            return 1e-7
        return LR
```

```
In [0]: #sgd=SGD(lr=0.05,decay=1e-5)
```

```
In [14]: model_vgg_scene.compile(loss='categorical_crossentropy',optimizer=optimizer)
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/optimizers.py:793: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3576: The name tf.log is deprecated. Please use tf.math.log instead.

```
In [15]: keras.callbacks.LearningRateScheduler(learning_rate_schedule)
```

```
Out[15]: <keras.callbacks.LearningRateScheduler at 0x7efe0201a780>
```

```
In [16]: history=model_vgg_scene.fit(train_x, train_y, validation_data=(test_x, test_y))
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow_core/python/ops/math_grad.py:1424: where (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:1033: The name tf.assign_add is deprecated. Please use tf.compat.v1.assign_add instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:1020: The name tf.assign is deprecated. Please use tf.compat.v1.assign instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3005: The name tf.Session is deprecated. Please use tf.compat.v1.Session instead.

Train on 4037 samples, validate on 448 samples

Epoch 1/30

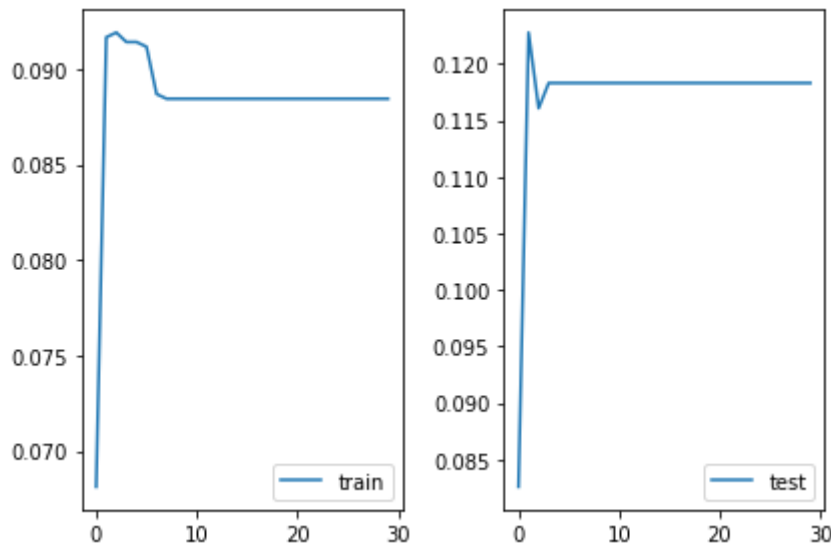
```
In [0]: import matplotlib.pyplot as plt
import numpy as np
from keras import optimizers
from keras.callbacks import EarlyStopping
```

```
In [18]: plt.subplot(1, 2, 1)
plt.plot(history.history['acc'])
plt.legend(['train'], loc='lower right')
plt.subplot(1, 2, 2)

plt.plot(history.history['val_acc'])
plt.legend(['test'], loc='lower right')

plt.tight_layout()

plt.show()
```



VGG16 Pre-trained

This part using Pre-trained model ,and I design two parts. First is run 1 times and second is run 3 times .it is clearly that perform one times the accuracy increase a lot, even perform three times the model drop in local-minium

Perform Training one times

By using VGG16 pre-trained model run one time ,the accuracy increased to 78.1%

```
In [19]: ishape=224
model_vgg = VGG16(include_top=False, weights='imagenet', input_shape=(ishap
for layer in model_vgg.layers:
    layer.trainable = False
```

Downloading data from https://github.com/fchollet/deep-learning-models/releases/download/v0.1/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5 (https://github.com/fchollet/deep-learning-models/releases/download/v0.1/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5)
58892288/58889256 [=====] - 3s 0us/step

```
In [0]: x = model_vgg.get_layer('block5_pool').output
x = Flatten(name='flatten')(x)
x = Dense(512, activation='relu', name='fc1')(x)
x = Dense(4096, activation='relu', name='fc2')(x)
x = Dense(15, activation='softmax', name='predictions')(x)
```

```
In [21]: #model=Dense(15,activation='softmax')(model)
model_vgg_pre_scene=Model(inputs=model_vgg.input, outputs=x)
model_vgg_pre_scene.summary()
```

Model: "model_2"

Layer (type)	Output Shape	Param #
=====		
input_2 (InputLayer)	(None, 224, 224, 3)	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
flatten (Flatten)	(None, 25088)	0
fc1 (Dense)	(None, 512)	12845568
fc2 (Dense)	(None, 4096)	2101248
predictions (Dense)	(None, 15)	61455
=====		
Total params: 29,722,959		
Trainable params: 15,008,271		

Non-trainable params: 14,714,688

```
In [0]: # model = Dense(15, activation="softmax")(model)
# model_vgg_scene_pretrain = Model(model_vgg.input, model, name='vgg_16_pre
# model_vgg_scene_pretrain.summary()
```

```
In [0]: def learning_rate_schedule(epoch):
    if epoch <= 10:
        return 1e-4 # 0.00001
    elif epoch <= 20:
        return 1e-5
    elif epoch <= 30:
        return 1e-6
    else:
        return 1e-7
    return LR
```

```
In [0]: #sgd = SGD(lr = 0.01, decay = 1e-5)
model_vgg_pre_scene.compile(loss='categorical_crossentropy', optimizer=opti
```

```
In [25]: keras.callbacks.LearningRateScheduler(learning_rate_schedule)
```

```
Out[25]: <keras.callbacks.LearningRateScheduler at 0x7efdf01c5e80>
```

```
In [26]: model_vgg_pre_scene.fit(train_x, train_y, validation_data=(test_x, test_y),
```

```
Train on 4037 samples, validate on 448 samples
Epoch 1/10
4037/4037 [=====] - 44s 11ms/step - loss: 1.0449
- acc: 0.6757 - val_loss: 0.5533 - val_acc: 0.7835
Epoch 2/10
4037/4037 [=====] - 43s 11ms/step - loss: 0.3380
- acc: 0.8922 - val_loss: 0.4819 - val_acc: 0.8348
Epoch 3/10
4037/4037 [=====] - 43s 11ms/step - loss: 0.1855
- acc: 0.9403 - val_loss: 0.4864 - val_acc: 0.8170
Epoch 4/10
4037/4037 [=====] - 43s 11ms/step - loss: 0.0925
- acc: 0.9755 - val_loss: 0.4586 - val_acc: 0.8482
Epoch 5/10
4037/4037 [=====] - 43s 11ms/step - loss: 0.0614
- acc: 0.9854 - val_loss: 0.4389 - val_acc: 0.8594
Epoch 6/10
4037/4037 [=====] - 43s 11ms/step - loss: 0.0324
- acc: 0.9916 - val_loss: 0.4209 - val_acc: 0.8750
Epoch 7/10
4037/4037 [=====] - 43s 11ms/step - loss: 0.0238
- acc: 0.9946 - val_loss: 0.4572 - val_acc: 0.8638
Epoch 8/10
4037/4037 [=====] - 43s 11ms/step - loss: 0.0155
- acc: 0.9983 - val_loss: 0.4562 - val_acc: 0.8683
Epoch 9/10
4037/4037 [=====] - 43s 11ms/step - loss: 0.0151
- acc: 0.9973 - val_loss: 0.4828 - val_acc: 0.8504
Epoch 10/10
4037/4037 [=====] - 43s 11ms/step - loss: 0.0103
- acc: 0.9990 - val_loss: 0.4548 - val_acc: 0.8616
```


```
Out[26]: <keras.callbacks.History at 0x7efdf018d6d8>
```

```
In [27]: plt.subplot(1, 2, 1)
plt.plot(model_vgg_pre_scene.history.history['acc'])
plt.legend(['train'], loc='lower right')
plt.subplot(1, 2, 2)

plt.plot(model_vgg_pre_scene.history.history['val_acc'])
plt.legend(['test'], loc='lower right')

plt.tight_layout()

plt.show()
```



Perform Training Three Times

The code which show below is not good way to normlization data in this dataset

```
In [0]: # from tensorflow.keras.applications.vgg16 import preprocess_input
# x_test = preprocess_input(test_x)
# x_train = preprocess_input(train_x)
```

```
In [29]: ishape=224

model = VGG16(weights='imagenet', include_top=False, input_shape=(ishape, ishape, 3))
for layer in model_vgg.layers:
    layer.trainable = False
print(model.summary())
```

Model: "vgg16"

Layer (type)	Output Shape	Param #
=====		
input_3 (InputLayer)	(None, 224, 224, 3)	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
=====		
Total params: 14,714,688		
Trainable params: 14,714,688		
Non-trainable params: 0		
None		

```
In [0]: from keras.layers import Flatten, Dense
        from keras.models import Model

        x = model.get_layer('block5_pool').output
        x = Flatten(name='flatten')(x)
        x = Dense(512, activation='relu', name='fc1')(x)
        #x = Dense(4096, activation='relu', name='fc2')(x)
        x = Dense(15, activation='softmax', name='predictions')(x)

        model_pre_updated = Model(inputs=model.input, outputs=x)
```

```
In [0]: def learning_rate_schedule(epoch):
        if epoch <= 10:
            return 1e-4 # 0.00001
        elif epoch <= 20:
            return 1e-5
        elif epoch <= 30:
            return 1e-6
        else:
            return 1e-7
        return LR
```

```
In [0]: model_pre_updated.save_weights('model_initial.h5')
```

```

In [33]: from keras import optimizers
from keras.callbacks import EarlyStopping

model_pre_updated.save_weights('model_initial.h5')
training_runs = []
for i in range(3):
    #model_updated.compile(loss='categorical_crossentropy', optimizer='sgd'

    model_pre_updated.compile(loss='categorical_crossentropy', optimizer=op
    keras.callbacks.LearningRateScheduler(learning_rate_schedule)
    history = model_pre_updated.fit(train_x, train_y, batch_size=32, shuff
    training_runs.append(history)
    model_pre_updated.get_weights()
    if i == 2:
        model_pre_updated.save_weights('model1.h5')
    else:
        model_pre_updated.load_weights('model_initial.h5')
print()

```

Train on 4037 samples, validate on 448 samples

Epoch 1/10

4037/4037 [=====] - 119s 29ms/step - loss: 1.025
2 - acc: 0.6554 - val_loss: 0.6536 - val_acc: 0.7656

Epoch 2/10

4037/4037 [=====] - 118s 29ms/step - loss: 0.425
8 - acc: 0.8556 - val_loss: 0.5316 - val_acc: 0.8058

Epoch 3/10

4037/4037 [=====] - 118s 29ms/step - loss: 0.279
4 - acc: 0.9009 - val_loss: 0.7661 - val_acc: 0.7545

Epoch 4/10

4037/4037 [=====] - 117s 29ms/step - loss: 0.211
7 - acc: 0.9289 - val_loss: 0.4589 - val_acc: 0.8393

Epoch 5/10

4037/4037 [=====] - 117s 29ms/step - loss: 0.132
0 - acc: 0.9562 - val_loss: 0.6158 - val_acc: 0.8237

Epoch 6/10

4037/4037 [=====] - 117s 29ms/step - loss: 0.095
8 - acc: 0.9693 - val_loss: 0.5353 - val_acc: 0.8348

Epoch 7/10

4037/4037 [=====] - 117s 29ms/step - loss: 0.078
9 - acc: 0.9710 - val_loss: 0.6736 - val_acc: 0.8192

Epoch 8/10

4037/4037 [=====] - 117s 29ms/step - loss: 0.072
4 - acc: 0.9757 - val_loss: 0.5754 - val_acc: 0.8371

Epoch 9/10

4037/4037 [=====] - 117s 29ms/step - loss: 0.074
6 - acc: 0.9720 - val_loss: 0.5155 - val_acc: 0.8594

Epoch 10/10

4037/4037 [=====] - 117s 29ms/step - loss: 0.068
7 - acc: 0.9780 - val_loss: 0.5576 - val_acc: 0.8527

Train on 4037 samples, validate on 448 samples

Epoch 1/10

4037/4037 [=====] - 119s 30ms/step - loss: 1.160
0 - acc: 0.6178 - val_loss: 0.8131 - val_acc: 0.7388

Epoch 2/10

4037/4037 [=====] - 118s 29ms/step - loss: 0.498

```
5 - acc: 0.8283 - val_loss: 0.5978 - val_acc: 0.7924
Epoch 3/10
4037/4037 [=====] - 118s 29ms/step - loss: 0.272
4 - acc: 0.9036 - val_loss: 0.3909 - val_acc: 0.8594
Epoch 4/10
4037/4037 [=====] - 117s 29ms/step - loss: 0.202
0 - acc: 0.9274 - val_loss: 0.5724 - val_acc: 0.8371
Epoch 5/10
4037/4037 [=====] - 117s 29ms/step - loss: 0.147
7 - acc: 0.9475 - val_loss: 0.4733 - val_acc: 0.8393
Epoch 6/10
4037/4037 [=====] - 117s 29ms/step - loss: 0.095
1 - acc: 0.9663 - val_loss: 0.5295 - val_acc: 0.8460
Epoch 7/10
4037/4037 [=====] - 117s 29ms/step - loss: 0.122
8 - acc: 0.9567 - val_loss: 0.4194 - val_acc: 0.8594
Epoch 8/10
4037/4037 [=====] - 117s 29ms/step - loss: 0.051
4 - acc: 0.9841 - val_loss: 0.4565 - val_acc: 0.8705
Epoch 9/10
4037/4037 [=====] - 117s 29ms/step - loss: 0.031
4 - acc: 0.9896 - val_loss: 0.6239 - val_acc: 0.8460
Epoch 10/10
4037/4037 [=====] - 117s 29ms/step - loss: 0.071
1 - acc: 0.9762 - val_loss: 0.7228 - val_acc: 0.8214
```

Train on 4037 samples, validate on 448 samples

```
Epoch 1/10
4037/4037 [=====] - 120s 30ms/step - loss: 1.094
0 - acc: 0.6289 - val_loss: 0.5747 - val_acc: 0.8147
Epoch 2/10
4037/4037 [=====] - 118s 29ms/step - loss: 0.477
1 - acc: 0.8400 - val_loss: 0.6016 - val_acc: 0.7835
Epoch 3/10
4037/4037 [=====] - 118s 29ms/step - loss: 0.306
1 - acc: 0.8960 - val_loss: 0.5267 - val_acc: 0.8103
Epoch 4/10
4037/4037 [=====] - 118s 29ms/step - loss: 0.166
6 - acc: 0.9453 - val_loss: 0.4529 - val_acc: 0.8683
Epoch 5/10
4037/4037 [=====] - 118s 29ms/step - loss: 0.160
6 - acc: 0.9472 - val_loss: 0.3962 - val_acc: 0.8683
Epoch 6/10
4037/4037 [=====] - 117s 29ms/step - loss: 0.088
6 - acc: 0.9700 - val_loss: 0.4773 - val_acc: 0.8594
Epoch 7/10
4037/4037 [=====] - 117s 29ms/step - loss: 0.119
2 - acc: 0.9596 - val_loss: 0.5349 - val_acc: 0.8393
Epoch 8/10
4037/4037 [=====] - 117s 29ms/step - loss: 0.095
5 - acc: 0.9678 - val_loss: 0.3961 - val_acc: 0.8705
Epoch 9/10
4037/4037 [=====] - 117s 29ms/step - loss: 0.047
7 - acc: 0.9819 - val_loss: 0.7749 - val_acc: 0.8281
Epoch 10/10
4037/4037 [=====] - 117s 29ms/step - loss: 0.038
```

1 - acc: 0.9876 - val_loss: 0.4519 - val_acc: 0.8817

```
In [0]: import matplotlib.pyplot as plt  
import numpy as np
```



```

In [35]: plt.subplot(2, 3, 1)
plt.plot(training_runs[0].history['acc'])
plt.legend(['train'], loc='lower right')

plt.subplot(2, 3, 2)
plt.plot(training_runs[1].history['acc'])
plt.legend(['train'], loc='lower right')

plt.subplot(2, 3, 3)
plt.plot(training_runs[2].history['acc'])
plt.legend(['train'], loc='lower right')

plt.subplot(2, 3, 4)
plt.plot(training_runs[0].history['val_acc'])
plt.legend(['test'], loc='lower right')

plt.subplot(2, 3, 5)
plt.plot(training_runs[1].history['val_acc'])
plt.legend(['test'], loc='lower right')

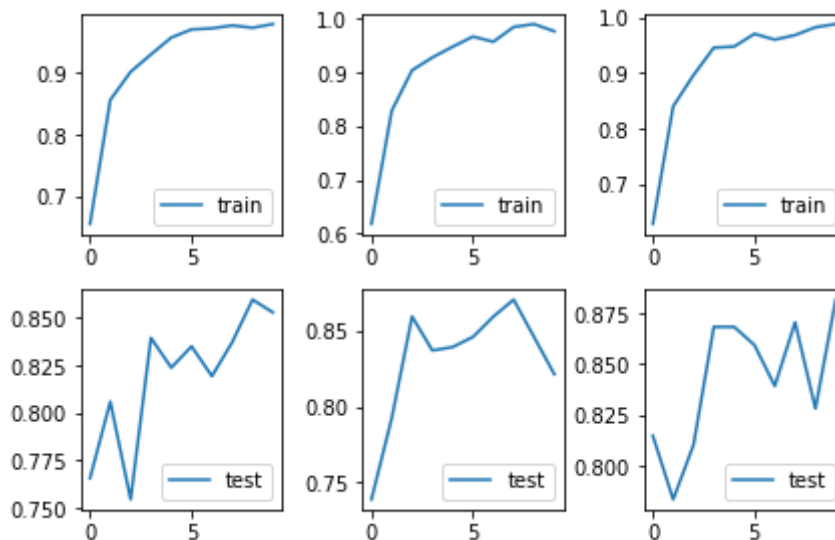
plt.subplot(2, 3, 6)
plt.plot(training_runs[2].history['val_acc'])
plt.legend(['test'], loc='lower right')

plt.tight_layout()

plt.show()

print("Average training accuracy: {}".format(np.mean([training_runs[0].hist
                                                    training_runs[1].histo
print("Average testing accuracy: {}".format(np.mean([training_runs[0].histo
                                                    training_runs[1].histo

```



Average training accuracy: 0.9805961522631992
Average testing accuracy: 0.8519345238095237

